

I already uploaded the report within the previous time constraints. But I wanted to upload a revised version since the deadline was postponed.

(a) Data Generation and CSV Export

Two 2D Gaussian classes (each with 500 samples) were generated and saved to `distribution.csv`. Parameters used:

$$\mu_1 = \begin{bmatrix} 1.5 \\ -0.5 \end{bmatrix}, \quad \mu_2 = \begin{bmatrix} -1.0 \\ 0.8 \end{bmatrix}, \quad \Sigma = \begin{bmatrix} 1.2 & -0.4 \\ -0.4 & 0.8 \end{bmatrix}.$$

Essentials:

```
n_samples = 500
mean1 = [1.5, -0.5]
mean2 = [-1.0, 0.8]
cov = [[1.2, -0.4], [-0.4, 0.8]]

X1 = np.random.multivariate_normal(mean1, cov, n_samples)
X2 = np.random.multivariate_normal(mean2, cov, n_samples)
y1 = np.ones(n_samples); y2 = -np.ones(n_samples)

X = np.vstack((X1, X2))
y = np.hstack((y1, y2))

df = pd.DataFrame({"x1":X[:,0], "x2":X[:,1], "label":y})
df.to_csv("distribution.csv", index=False)
```




Figure 1: Scatter plot of the generated data

(b) Perceptron Learning Algorithm

Update rule (only if $y_i(w^\top x_i + b) \leq 0$):

The method given in the course files was updated by shuffling the data in each epoch. The reason for this was although the loss value was reasonable in the loss function, the classification boundary line was graphed in a way that it classified the entire data as either one of the classes. I depicted the reason for this was the loss function was optimized according to a local minima in the gradient descent process. And since the data was small enough to do these kind of optimizations in the stochastic gradient descent process, this method was used to prevent it.

10, 20, 50 and 100 epochs were tested in this code to analyze the effects of different epoch numbers on the optimization. The learning rate was chosen 0.5 after some testing to prevent small epoch numbers from performing too poor for comparison.

```
def perceptron_train(X, y, lr=0.05, epochs=20, seed=42):
    rng = np.random.default_rng(seed)
    n, d = X.shape
    w = np.random.rand(d); b = np.random.rand()
    for epoch in range(epochs):
        idx = rng.permutation(n) # shuffle each epoch
        for i in idx:
            margin = y[i] * (np.dot(w, X[i]) + b)
            if margin <= 0:
                w += lr * y[i] * X[i]
                b += lr * y[i]
    return w, b

# Start training
w_0, b_0, fl_0 = perceptron_train(X, y, lr=0.05, epochs=10)
w_1, b_1, fl_1 = perceptron_train(X, y, lr=0.05, epochs=20)
w_2, b_2, fl_2 = perceptron_train(X, y, lr=0.05, epochs=50)
w_3, b_3, fl_3 = perceptron_train(X, y, lr=0.05, epochs=100)
```

(c) Loss vs. Epoch (10, 20, 50, 100)

Average misclassification proxy per epoch recorded for 4 runs (10/20/50/100 epochs).

Figure 2. Loss progression subplots for 4 epoch counts.




Figure 2: Loss progression plots

(d) Decision Boundary

Decision boundary:

$$w_1x_1 + w_2x_2 + b = 0 \Rightarrow x_2 = -\frac{w_1}{w_2}x_1 - \frac{b}{w_2}.$$

Numerical equations obtained from trained weights/biases:

- **Epochs 10:** $w = [0.1052, -0.0750]$, $b = 0.0533 \Rightarrow x_2 = 1.401x_1 + 0.710$
- **Epochs 20:** $w = [0.1440, -0.1176]$, $b = -0.1102 \Rightarrow x_2 = 1.225x_1 + 0.937$
- **Epochs 50:** $w = [0.0748, 0.0198]$, $b = 0.0347 \Rightarrow x_2 = -3.786x_1 - 1.756$
- **Epochs 100:** $w = [0.0308, -0.1410]$, $b = -0.0351 \Rightarrow x_2 = 0.219x_1 - 0.249$

Summary of Learned Parameters and Metrics

Epochs	Slope	Intercept	Acc	Prec	F1
10	1.40	0.71	0.86	0.80	0.87
20	1.22	0.94	0.87	0.93	0.86
50	-3.79	-1.76	0.81	0.74	0.84
100	0.22	-0.25	0.81	0.88	0.80




Figure 3: Decision boundary plots

(e) Sample Means and Covariances (by Class)

Class +1 ($y=+1$):

$$\hat{\mu}_1 = \begin{bmatrix} 1.5571 \\ -0.5572 \end{bmatrix}, \quad \hat{\Sigma}_1 = \begin{bmatrix} 1.1558 & -0.3629 \\ -0.3629 & 0.7724 \end{bmatrix}.$$

Class -1 ($y=-1$):

$$\hat{\mu}_2 = \begin{bmatrix} -1.0203 \\ 0.8172 \end{bmatrix}, \quad \hat{\Sigma}_2 = \begin{bmatrix} 1.1502 & -0.3775 \\ -0.3775 & 0.7371 \end{bmatrix}.$$

(f) Confusion Matrices

Layout:

$$\begin{bmatrix} TP & FN \\ FP & TN \end{bmatrix}$$

- Epochs 10:** $\begin{bmatrix} 470 & 30 \\ 90 & 410 \end{bmatrix}$, Acc=0.88, Prec=0.84, Rec=0.94, F1=0.89.
- Epochs 20:** $\begin{bmatrix} 477 & 23 \\ 119 & 381 \end{bmatrix}$, Acc=0.86, Prec=0.80, Rec=0.95, F1=0.87.
- Epochs 50:** $\begin{bmatrix} 459 & 41 \\ 135 & 365 \end{bmatrix}$, Acc=0.82, Prec=0.77, Rec=0.92, F1=0.84.
- Epochs 100:** $\begin{bmatrix} 415 & 85 \\ 37 & 463 \end{bmatrix}$, Acc=0.88, Prec=0.92, Rec=0.83, F1=0.87.




Figure 4: Confusion matrix plots for each epoch

(g) Performance Metrics (20 Epochs)

For the 10-epoch model: Accuracy=0.88, Precision=0.84, Recall=0.94, $F_1=0.89$.
 For the 20-epoch model: Accuracy=0.87, Precision=0.93, Recall=0.79, $F_1=0.86$.
 For the 50-epoch model: Accuracy=0.82, Precision=0.77, Recall=0.92, $F_1=0.84$.
 For the 100-epoch model: Accuracy=0.88, Precision=0.92, Recall=0.83, $F_1=0.87$.

(h) Discussion and Efficiency Evaluation

The perceptron achieved the best trade-off at **20 epochs** with accuracy 0.87 and F1 score 0.86. The slope and intercept changes over training show the boundary adjusting directionally as the algorithm converges. My deduction from this result is that lower epoch numbers may cause the optimization end prematurely and the higher epoch rates can cause oscillations in the loss rate. Whether or not the loss rate being equal to the lower values of this oscillations or to the higher ones cannot be predicted simply, so keeping the epoch number in a reasonable rate actually improves the converged loss rate.

The model correctly captures the linear separability trend between the two Gaussian clusters. However, because the Gaussians overlap, 100% accuracy is impossible; errors concentrate in the intersection region.

Considering the theoretical Bayes error (around 10–15% for such overlap), achieving 81–87% accuracy indicates the model performs **close to optimally for a linear classifier**. It learns effectively, showing that the algorithm is **efficient enough and behaves correctly**, even though it will not outperform more advanced probabilistic or kernel-based models.

Bayes Error Explanation. For any two probabilistic classes $P(x|C_1)$ and $P(x|C_2)$ with prior probabilities $P(C_1)$ and $P(C_2)$, the minimum possible classification error is given by the *Bayes error rate*:

$$P_e = \int \min [P(C_1|x), P(C_2|x)] dx.$$

This represents the theoretical limit of how well any classifier—even an ideal one—can separate the two distributions. Because the Gaussian distributions in this experiment overlap significantly, the Bayes error is nonzero (approximately 10–15%). Therefore, the perceptron’s achieved accuracy (up to 87%) is very close to the theoretical maximum achievable for this dataset, demonstrating that it is indeed a reasonable and efficient model for linearly separable approximation. Although it is still better to evaluate this result relative to the task which is tried to be accomplished. For example an accuracy rate of 87% is not reasonable in the medical field, but it may be accepted in a spam mail detection task.

The handwritten solutions are included at the end of the report. The [GitHub Link](#) and [CoLab](#) link for this homework are in the underlined relative texts.

Ata Ataş

040230771

Ataş

b) Dağılımlar arası örtüşme sağlayabilmek için mean₁ = $\begin{bmatrix} 2,5 & -0,5 \end{bmatrix}$ ~~mean1~~
mean₂ = $\begin{bmatrix} -4,0 & 0,8 \end{bmatrix}$

ve varyansları da var₁ = 1,2; var₂ = 0,8 şeklinde farklı seviyelerde optimizasyon için daha karmaşık bir ortam yaratmaktadır.

Kovaryansı -0,4 olmak üzere iki sınıf arası negatif bir ilişki yaratmaktadır.

$$\text{cov} = \begin{bmatrix} 1,2 & -0,4 \\ -0,4 & 0,8 \end{bmatrix}$$

c) Epoch sayısı 10'dan 20'ye giderek loss azalıyor. Bunun sebebi, daha fazla iterasyon sayesinde gradient minimum noktasına yakının daha kolaylaşması, olacak verimlilik. Ama 20 Epoch içinde loss'un artması, da momentum rate ile fazla iterasyon sayısının gradient hesabında oscilasyonlara sebebi olması olacak açıklayabiliriz.

c) class 1 mean = $\begin{bmatrix} 1,56 & -0,56 \end{bmatrix}$ cov: $\begin{bmatrix} 4,58 & -0,36 \\ -0,36 & 0,77 \end{bmatrix}$

class 2 mean = $\begin{bmatrix} -1,020 & 0,82 \end{bmatrix}$ cov: $\begin{bmatrix} 1,15 & -0,37 \\ -0,37 & 0,75 \end{bmatrix}$

h) Standart bazes öğrenicinde maksimum verimlilik: %86'ya yetkin bulunken %87 perceptron accuracy'yu gerçekleştirmek için her bir gauss dağılımına göre işi sayılır.